

Towards Dynamic Switching in Per-Run Algorithm Selection

July 9, 2025




Continuous Black-Box Optimization

- ▶ Goal: Optimize a function $f : \mathbb{R}^D \rightarrow \mathbb{R}$
- ▶ Black-Box: Optimization is limited to sampling

Per-Instance Algorithm Selection


1. Sample f randomly
2. Compute ELA features from samples
3. Use features as input for a machine learning model that chooses the algorithm
4. Run the algorithm

Per-Run Algorithm Selection



- ▶ Approach from Kostovska et al. 2022:
 1. After 150 evaluations, calculate ELA features from CMA-ES samples
 2. Random Forest regression models predict target precisions of the six A2 algorithms
 3. Choose A2 algorithm with lowest predicted precision
 4. Warm-start second algorithm
- ▶ A1 budget is static!

Why is Dynamic Switching Necessary?




- ▶ Recording of 2400 runs conducted on BBOB functions

BBOB Test Suite (Hansen et al. 2016)


- ▶ Set of 24 synthetic noiseless black-box functions
- ▶ Each function has several instances
- ▶ Standard benchmarking set

Why is Dynamic Switching Necessary?




General Setup

1. Switching Decision



2. Algorithm Selection



First Approach

- ▶ For each run, we try to find the switching points defined as above
- ▶ Selection process consists of two parts:

Switching Decision

- ▶ Let s_1, \dots, s_n be the considered switching points
- ▶ At switching point s_i , we predict the precision of the best algorithm at s_i, s_{i+1}, \dots, s_n
- ▶ We switch if the predicted precision of the current switching point is the lowest

Algorithm Selection


- ▶ Predict the precisions of all six algorithms, switch to the algorithm with the lowest precision

First Approach: Our setup

- ▶ Switching points: 8, 16, ..., 96, 100, 150, 200, ..., 1000
- ▶ Machine learning input:
 - ▶ ELA features from CMA-ES samples
 - ▶ CMA-ES internal state during the last iteration
- ▶ Model training on first 5 instances of all BBOB functions, 20 runs each
- ▶ Evaluation on instances 6 and 7, 20 runs each
- ▶ Parameter tuning using ASF
- ▶ Baselines:
 - ▶ SBS: (Algorithm, Budget) combination that is best on the training set
 - ▶ VBS: Chooses the best switching point and algorithm at that switching point for each run
- ▶ Metric:

$$cg_{selector} = \frac{m_{SBS} - m_{selector}}{m_{SBS} - m_{VBS}}$$

First Approach: Results



Poor Performance

- ▶ Switching decision based on performance of best algorithm
⇒ Switch too early
- ▶ ELA features do not capture enough run-specific information

Second Approach

Let s_1, \dots, s_n be the switching points, a_1, \dots, a_n be the algorithm selectors and sd_1, \dots, sd_n be the switching decision models

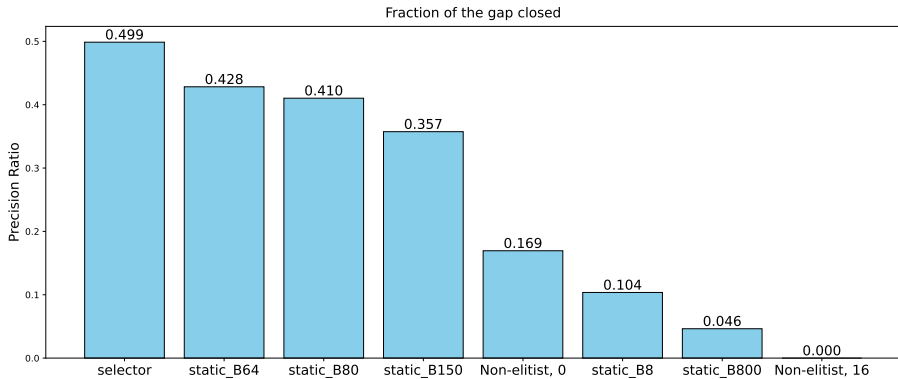
1. Evaluate Performances of a_1, \dots, a_n

1. Record the performances of a_1, \dots, a_n on first five instances
2. For each function, determine which a_i performed best
 $\Rightarrow s_i$ is optimal for that function


2. Training data for sd_1, \dots, sd_n

1. For ELA-features and CMA-ES data belonging to a run at s_i , we label them as true iff s_i is greater or equal than the optimal switching point of the run's function
2. For each s_i , train a binary classifier, predicting whether or not to switch

Second Approach: Results



Second Approach: Results



Permutation Test: Selector vs Static Baseline

Setup

- ▶ Data: Precision values per run for selector and static baseline selector
- ▶ Test: Permutation test using SciPy with 10,000 resamples
- ▶ Statistic: Mean difference in reached precision

Null Hypothesis

- ▶ Null Hypothesis (H_0):
The selector and static baseline have the same distribution of precision values. Any observed difference is due to random chance.

⇒ Found statistical significance


Next Step: In-Depth Empirical Evaluation

- ▶ Repeat experiments for dimension 10
- ▶ BBOB functions are purely synthetic
- ▶ Evaluation of the selector on LassoBench (Šehić et al. 2022)
 - ▶ Real-world HPO functions for weighted lasso regression
- ▶ Function from LassoBench do not have instances
⇒ Train-test splits using different runs on each function
- ▶ And other real-world benchmarks


Take-Home Message

- ▶ Dynamic Switching leads to improvements over static switching
- ▶ First approach did not work
 - ▶ ELA features do not seem to capture run-specific information
 - ▶ Suboptimal performance of algorithm selectors across budgets
- ▶ Second approach leads to significant improvements over static switching
- ▶ Next step: Evaluate robustness of the selector

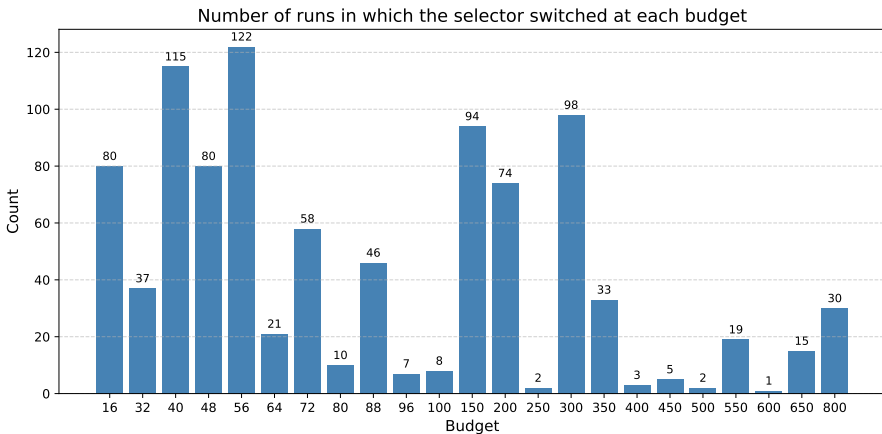
Permutation Test Results



Switching Budgets: First Approach



Switching Budgets: Second Approach



Algorithm Portfolio (Schröder et al. 2022)

- ▶ Broyden-Fletcher-Goldfarb-Shanno (BFGS): Deterministic line-search, approximates the hessian, local search
- ▶ Differential Evolution (DE): Random mutation of target solutions using binomial crossover
- ▶ Multi-Level Single Linkage (MLSL): Local searches based on clustering heuristics
- ▶ Particle Swarm Optimization (PSO): Swarm-based, velocity of particles gets adjusted based on their own and the global optimum
- ▶ CMA-ES elitism/non-elitism: In each iteration, does the offspring become the new population or does the new population consist of the best samples from both the old population and the offspring?