

Week 15

- I re-run the experiments with and without normalization of the data, using I-BFGS-b or BFGS with clipping
- Normalization+Switching Model Tuning worked really well this time, both outperform all static switching models on the test set, in case of BFGS with clipping also with statistical significance
- Results in the repository under “results week 15”
- I also tried optimizing the decision thresholds but this didn't work
- Performance improvements for both clipped BFGS and I_BFGS_b, statistical significance only for clipped BFGS. This is probably because for I_BFGS_b, some static switching budgets are also really good

Questions

- Which BFGS version do we use?
- Next steps?