

Week 9: Function specific switching

- Idea: Measure the performance of the static algorithm selectors on instances 1 to 5
- I used a 4-fold cross validation where I trained the selectors on $\frac{3}{4}$ of runs, and then let them predict the best algorithm for the remaining quarter.
- For each function, I set the switching point as the budget for which the static selector performed at that budget best on that function (lowest sum of precisions)
- I then trained random forests, one for each possible switching budget, that predict whether or not to switch for a run (so a binary prediction here)

Week 9: Function specific switching

- In a first attempt, I set the switch decision to true for a function and a budget, only if that budget is the right budget for the function
- This resulted in poor performance, because each run for which the right budget was not detected, then continued without switching, which is basically always a bad decision
- So, I set the switch decision to true for a function and a budget, if that budget is greater or equal than the right budget for that function. (So that if a run is not detected to switch at the right budget, it will switch somewhere behind that)
- This resulted in good performance

Results if we consider multiples of 50 as switching points

selector_optimal_switching:
Precision of the selector that,
for each fid, switches at the
best static switching point
for that fid (according to the
algorithm predictions on the
training runs)

Method	Ratio
selector	0.6914739331516234
selector_optimal_switching	0.6937242813091179
static_B50	0.7342292114590188
static_B200	0.7834688934562055
static_B100	0.8096604126452549
static_B150	0.8279748194510756
static_B350	0.8295645201063693
static_B450	0.8315831727495465
static_B300	0.833793375248887
static_B400	0.8454845500999881
static_B250	0.8521377125549351
static_B500	0.8646264151484486
static_B550	0.8651682121985461
static_B600	0.8680715726930405
static_B650	0.8727674677814952
static_B700	0.9691203768199774
static_B750	1.170709233204784
static_B800	1.2944068941955356
static_B850	1.612240193937837
static_B900	1.9356077011386117
static_B950	2.5804227966401223
static_B1000	3.9009662685339572

Results for multiples of 50 and multiples of 8 lower than 100

Method	Ratio
selector	0.7348407859659998
selector_optimal_switching	0.7414546343743612
static_B80	0.8343698876143243
static_B56	0.8431062571644096
static_B64	0.865747418437369
static_B24	0.8695766953339663
static_B40	0.874061958394088
static_B96	0.8749348244632318
static_B200	0.8864525814801185
static_B100	0.9095093121298923
static_B48	0.9111641348804176
static_B150	0.9256317196040164
static_B350	0.9270311535031466
static_B450	0.9288081993450942
static_B32	0.9297976509890551
static_B72	0.9301458725592493
static_B300	0.9307538689849849
static_B400	0.9410457603249096
static_B250	0.9469026246627765
static_B88	0.9556409676800052
static_B500	0.9578965899148633
static_B550	0.9583735408149151
static_B600	0.9609294063383146
static_B650	0.9650632630815067
static_B16	0.9933172712163364
static_B700	1.049883966110878
static_B750	1.227345224853125
static_B800	1.3362378641416175
static_B8	1.3913457921804038
static_B850	1.6160305989880546
static_B900	1.9006951677212154
static_B950	2.4683341759274477
static_B1000	3.63082555044471

Evaluation on new instances

- I also wanted to see if this works if we use runs on instances 6 and 7 (20 each) as a test set. Once again, I considered both sets of switching points
- The performance is worse here. Maybe this is because the distributions of training and test set are too different?

Method	Ratio
selector_precision	0.7397463435645291
static_B64	0.783110528315716
static_B80	0.7991784183080536
static_B96	0.8011234209716904
static_B56	0.8052528171186433
static_B48	0.8091478061854296
static_B100	0.8340053853588267
static_B72	0.8350493625064949
static_B550	0.8507966151371419
static_B300	0.8522574095492805
static_B400	0.8527850943443569
static_B350	0.857926650631531
static_B250	0.8623110386200014
static_B500	0.8793981739119423
static_B150	0.881782488942987
static_B600	0.8849761676564104
static_B450	0.8905819890916795
static_B650	0.8915562283902
static_B200	0.900379391354849
static_B88	0.9070488066040253
static_B24	0.9505310751659515
static_B700	0.9615915740558716
static_B32	0.9655715603363839
static_B16	1.011701664513621
static_B40	1.0258049077785734
static_B750	1.1243501419530186
static_B8	1.1955308516283178
static_B800	1.2689075171966206
static_B850	1.6472226745221685
static_B900	1.9641426220843177
static_B950	3.2024631000855757
static_B1000	8.075054376573444

Method	Ratio
selector_optimal_switching	0.7997233486711641
static_B50	0.8052528171186433
selector	0.8082308838341555
static_B100	0.8340053853588267
static_B550	0.8507966151371419
static_B300	0.8522574095492805
static_B400	0.8527850943443569
static_B350	0.857926650631531
static_B250	0.8623110386200014
static_B500	0.8793981739119423
static_B150	0.881782488942987
static_B600	0.8849761676564104
static_B450	0.8905819890916795
static_B650	0.8915562283902
static_B200	0.900379391354849
static_B700	0.9615915740558716
static_B750	1.1243501419530186
static_B800	1.2689075171966206
static_B850	1.6472226745221685
static_B900	1.9641426220843177
static_B950	3.2024631000855757
static_B1000	8.075054376573444

Questions

- Is the switching selection done right? I think we could probably improve here. Currently, around 80% of runs switch at the right switching point. We need a way to handle the remaining runs.
- I also tried to tune the switching models (according to F1 score), but this did not improve the performance. Should I tune them according to a different metric?
- The performance of the static selectors on the training and test data often does not really match. I think we need more data (so more runs per instance) to really understand this. Currently, the performance of a static depends on only a few (un)lucky algorithm choices